

Jaden, I need your honest read on what V3 actually proves.

The model is trained and frozen. My problem now is deciding what claim is fair.

Jaden Palmer

ARTISANS Lab, North Carolina State University

Your advice during our last call honestly changed the project. I finished the LSTM run, froze it before outside testing, and then tested it much harder than I originally planned. Some results are strong. Some are uncomfortable. I want the final project to show both.

THE SCORECARD I WOULD SHOW A JUDGE

RESULT	TEST	WHAT IT MEANS
0.97%	Matched synthetic test	Worst-mesh median Ac-225 endpoint error.
3.94%	Wider 10,000-case grid	Productive-case median; the p95 rose to 183.84%.
Rank 2	Schedule screening	V3 selected Radau's second-ranked schedule with 0.0470% objective regret.
Mixed	LSTM versus matched MLP	The evidence does not support saying the LSTM is always better.
Failed	Berkeley outside test	The two cases were outside training support and about 600 times too high.
Slower	Exact analytic baseline	The exact solver beats V3 on this reduced constant-rate system.

What I think I can defend

- V3 can be studied as a frozen surrogate for solver-verified screening inside a declared reduced domain.
- Its median ranking behavior can still be useful even when every finalist is checked by a trusted solver.
- The failure map and warning system are part of the result, not something to hide.

What I am still unsure about

- Which error statistic should be the headline without making the model look better than it is.
- How to define uncertainty and an out-of-range warning that a judge will accept.
- Whether the current model still has a useful research role when the exact reduced solver is faster.

The decision I need help making

If you were judging this project, what exact research question and claim would you accept from these results? I also need your opinion on the minimum additional evidence that matters most before December, especially if I cannot afford another full training run.